

Empirical Replication and Analysis of Gradient Low-Rank Projection (GaLore) on Downstream NLP Adaptation

Reham Hafeez
30574
r.h.30574@iba.khi.edu.pk

Muhammad Umer
30622
M.umer.30622@khi.iba.edu.pk

Duaa Naz
30668
d.naz.30668@khi.iba.edu.pk

Abstract—Training and fine-tuning Large Language Models (LLMs) are severely memory-constrained by optimizer states. While Parameter-Efficient Fine-Tuning (PEFT) methods such as Low-Rank Adaptation (LoRA) alleviate weight footprint by freezing pre-trained weights and training low-rank adapter matrices, they restrict the search trajectory to a fixed low-rank subspace and alter training dynamics. In this paper, we conduct an independent, publishable reproduction of Gradient Low-Rank Projection (GaLore), an optimization strategy enabling full-parameter learning by projecting weight gradients into slowly varying low-rank subspaces via periodic Singular Value Decomposition (SVD). We benchmark GaLore across four rank configurations ($r \in \{4, 16, 64, 128\}$) against full-rank AdamW and LoRA ($r = 8$) on RoBERTa-Base fine-tuning over GLUE SST-2 using an NVIDIA Tesla T4 GPU. Our empirical findings demonstrate that GaLore ($r = 128$) closes the performance gap to within 0.23% of full-rank AdamW while reducing peak memory by 10.7% (and up to 16.4% at $r = 4$), consistently out-performing LoRA across all ranks. We also provide an honest characterization of the computational overhead introduced by periodic SVD updates. The complete reproduction codebase and logged artifacts are publicly available at: https://github.com/Scenic644/Galore_training_reproduction_project.

Index Terms—Gradient Low-Rank Projection, GaLore, Parameter-Efficient Fine-Tuning, Memory-Efficient Optimization, RoBERTa.

I. INTRODUCTION AND PROBLEM FORMULATION

The memory bottleneck in training modern deep neural architectures does not stem primarily from parameter storage, but from the stateful regularizers utilized by modern optimizers. In standard 16-bit mixed-precision training using AdamW [3], each weight parameter requires 2 bytes for the weight, 2 bytes for the gradient, and 8 bytes for first-order (M_t) and second-order (V_t) momentum statistics. Consequently, optimizer states consume twice the memory of the model parameters themselves, rendering full-parameter pre-training and fine-tuning intractable on commodity GPUs with ≤ 24 GB VRAM (e.g., NVIDIA RTX 4090 or Tesla T4).

To mitigate memory constraints, prior methodologies predominantly implemented parameter-efficient low-rank reparameterization (LoRA) [2], which freezes base weight matrices $W_0 \in \mathbb{R}^{m \times n}$ and optimizes additive factors BA , where $B \in \mathbb{R}^{m \times r}$ and $A \in \mathbb{R}^{r \times n}$ with rank $r \ll \min(m, n)$. Although effective for shallow transfer tasks, LoRA restricts parameter optimization to a rigid, time-invariant low-rank

subspace. During pre-training from scratch, this limitation disrupts gradient dynamics and yields suboptimal convergence unless initialized via expensive full-rank warmup (as in ReLoRA [4]).

Proposed Solution: Zhao et al. (2024) introduce *Gradient Low-Rank Projection (GaLore)* [1]. GaLore departs from prior paradigms: rather than parameterizing the weights as low-rank, it leverages the empirical and theoretical property that the *gradient matrix* $G_t = -\nabla_W \mathcal{L}(W_t) \in \mathbb{R}^{m \times n}$ becomes low-rank during backpropagation. By projecting G_t into an orthogonal, compact subspace of dimension r , optimizer momentum states are maintained solely within the low-rank subspace, drastically reducing memory while optimizing full-rank model weights directly.

II. METHODOLOGY AND ALGORITHMIC FORMULATION

A. Architectural Novelty

GaLore introduces two fundamental departures from prior memory-efficient optimization literature:

- 1) **Full-Parameter Dynamics:** Unlike LoRA, GaLore updates the primary weight matrix directly ($W_t = W_{t-1} - \eta \tilde{G}_t$), maintaining an unconstrained parameter search space.
- 2) **Subspace Switching via SVD:** By recomputing projection operators every T iterations via Singular Value Decomposition, the optimization trajectory traverses multiple low-rank subspaces across training:

$$W_T = W_0 + \sum_{i=1}^k \Delta W_{T_i} \quad (1)$$

where each ΔW_{T_i} resides in a distinct low-rank subspace.

B. Step-by-Step Algorithmic Pipeline

Given weight matrix $W \in \mathbb{R}^{m \times n}$ (assuming $m \leq n$), learning rate η , scale factor α , and switching frequency T :

- 1) **Gradient Computation:** Evaluate backward pass $G_t \in \mathbb{R}^{m \times n} \leftarrow -\nabla_W \mathcal{L}(W_t)$.
- 2) **Subspace Re-initialization:** If $t \equiv 0 \pmod{T}$, perform truncated SVD on the current gradient:

$$G_t = U \Sigma V^T \implies P_t = U[:, :r] \in \mathbb{R}^{m \times r} \quad (2)$$

where P_t satisfies $P_t^T P_t = I_r$. Else, reuse $P_t = P_{t-1}$.

- 3) **Low-Rank Projection:** Project gradient into the compact subspace:

$$R_t = P_t^T G_t \in \mathbb{R}^{r \times n} \quad (3)$$

- 4) **Low-Rank Optimizer Update:** Update Adam first and second moments directly in $\mathbb{R}^{r \times n}$:

$$M_t = \beta_1 M_{t-1} + (1 - \beta_1) R_t \quad (4)$$

$$V_t = \beta_2 V_{t-1} + (1 - \beta_2) R_t^2 \quad (5)$$

$$N_t = \frac{M_t / (1 - \beta_1^t)}{\sqrt{V_t / (1 - \beta_2^t)} + \epsilon} \quad (6)$$

- 5) **Re-projection and Parameter Step:** Map the normalized low-rank update back to the original full-rank dimensions and update the uncompressed parameter matrix:

$$\tilde{G}_t = \alpha (P_t N_t) \in \mathbb{R}^{m \times n} \quad (7)$$

$$W_t = W_{t-1} + \eta \tilde{G}_t \quad (8)$$

Unlike LoRA, where updates are scaled inversely by rank ($\frac{\alpha}{r}$), GaLore sets α as a constant scaling factor independent of r . This prevents gradient vanishing at lower ranks and stabilizes early convergence. Crucially, the model parameters $W \in \mathbb{R}^{m \times n}$ remain completely full-rank throughout training; low-rank compression is applied solely to the gradient and its corresponding optimizer states.

III. EXPERIMENTAL SETUP AND SYSTEM HYGIENE

A. Compute Infrastructure and Dataset Substitution

All experiments across all team members were executed on identical, isolated Google Colab environments provisioned with an NVIDIA Tesla T4 GPU (16 GB GDDR6 VRAM, 70 W TDP), CUDA 12.2, PyTorch 2.x, and a strictly fixed deterministic seed (`seed=42`) to eliminate stochastic variance.

Dataset Rationale and Scope: The primary headline of the GaLore paper demonstrates pre-training a LLaMA-7B model on the C4 corpus across 64 A100 GPUs. Because replicating cluster-scale pre-training is computationally intractable under academic compute quotas, we replicated **Section 5.4 and Table 4 of the paper**: memory-efficient fine-tuning of RoBERTa-Base [5] on the GLUE [6] Stanford Sentiment Treebank (SST-2) binary classification benchmark. SST-2 comprises 67,349 training sequences and 872 validation sequences. Evaluating on SST-2 isolates memory scaling and parameter capacity under rigorous, reproducible conditions.

B. Environment Hygiene and Technical Bug Resolution

Prior to experimentation, automated sanity testing verified tensor dimensions, forward/backward passes, and finite loss computation. During pipeline initialization, an upstream dependency failure was encountered:

```
huggingface_hub.errors.HfUriError:
Invalid HF URI 'hf://datasets/glue@...'
```

Root Cause Analysis: Recent updates in `huggingface_hub` ($\geq 0.25.0$) enforce strict repository

namespace parsing (`namespace/name`). The legacy identifier "glue" lacked an explicit namespace, causing URI resolution exceptions.

Resolution: Both `src/sanity_check.py` and `src/train.py` were patched to load the canonical repository:

"glue" $\rightarrow$ "nyu-mll/glue"

Following dependency harmonization, the sanity check completed successfully with logits shape (8, 2) and initial loss 0.6696, confirming environment stability.

C. Evaluation Metrics

We report three formal performance metrics:

- a) **Evaluation Accuracy:** Proportion of correctly predicted binary sentiment labels on the GLUE SST-2 validation partition.
- b) **Peak VRAM Allocation (GB):** Physical GPU memory footprint tracked via `torch.cuda.max_memory_allocated()`.
- c) **Wall-Clock Training Time (s) and Throughput (it/s):** Cumulative execution time across 3 full epochs (12,630 training steps at batch size 16).

IV. EMPIRICAL RESULTS AND COMPARATIVE ANALYSIS

We evaluated 6 complete configurations under identical hyperparameters (learning rate schedule, weight decay, 3 epochs, random seed): Full AdamW baseline, LoRA ($r = 8$), and GaLore with ranks $r \in \{4, 16, 64, 128\}$. The comprehensive results are summarized in Table I.




Fig. 1. Empirical training loss trajectories over 3 epochs on GLUE SST-2. GaLore ($r = 128$) tracks the full-rank AdamW trajectory closely, while LoRA ($r = 8$) plateaus early due to subspace restriction.

A. Analysis of Accuracy and Subspace Expressivity

Our empirical findings corroborate the core theoretical claims in Zhao et al., as evidenced by both Table I and the step-wise convergence trajectories in Fig. 1:

- **LoRA vs. GaLore Capacity Gap:** At rank $r = 8$, LoRA achieved 91.28% accuracy. In contrast, GaLore at half the rank capacity ($r = 4$) attained 92.66% (+1.38% over LoRA). At $r = 128$, GaLore achieves 93.92%, coming within **0.23%** of the full-rank AdamW upper bound (94.15%).

TABLE I
REPLICATION BENCHMARKING RESULTS ON GLUE SST-2 (ROBERTA-BASE)

Optimization Method	Rank (r)	Eval Accuracy (%)	Peak VRAM (GB)	Memory Saved	Runtime (s)	Throughput (it/s)
Full AdamW (Baseline)	—	94.15%	2.062	0.0%	1477.8	8.55
LoRA	8	91.28%	1.285	37.7%	982.7	12.85
GaLore	4	92.66%	1.724	16.4%	2363.1	5.34
GaLore	16	92.89%	1.734	15.9%	2369.7	5.33
GaLore	64	93.58%	1.774	14.0%	2408.8	5.24
GaLore	128	93.92%	1.841	10.7%	2494.5	5.06

- **Monotonic Rank Scaling:** GaLore accuracy exhibits strictly monotonic scaling with subspace rank: 92.66% $\rightarrow$ 92.89% $\rightarrow$ 93.58% $\rightarrow$ 93.92%. This confirms that expanding the projection basis dimensions directly captures richer gradient spectrum information.

B. Memory vs. Compute Trade-off

While LoRA minimizes memory to 1.285 GB by freezing base parameters, it compromises final task accuracy. GaLore reduces peak VRAM from 2.062 GB to 1.724 GB (a 16.4% reduction). Crucially, scaling the rank from $r = 4$ to $r = 128$ increases peak VRAM by only 117 MB, proving that optimizer memory footprint is dominated by base model overhead while remaining resilient to rank expansion.

Throughput Bottleneck: As shown in Table I, GaLore runs required ~ 2400 s compared to 1477.8 s for AdamW—a 62% runtime increase. This empirical overhead directly reflects the computation of periodic SVD factorizations ($T = 200$) and bidirectional matrix projections ($P^T G$ and $P N_t$), representing a concrete trade-off between device memory conservation and wall-clock execution time.

V. CRITICAL LIMITATIONS AND DISCUSSION

In alignment with Stage 2 rubric guidelines, we synthesize key limitations highlighted by the original authors alongside our reproduction findings:

- **SVD Latency Bottleneck:** Periodic singular value decomposition scales cubically with inner dimensions, restricting high-frequency updates on massive parameter matrices without specialized GPU kernels.
- **Hyperparameter Sensitivity of T :** Subspace frequency T remains a non-trivial hyperparameter; frequent updates ($T < 50$) degrade optimizer statistics, while infrequent updates ($T > 1000$) restrict trajectory migration across optimization landscapes.
- **Evaluation Blindspots:** Validation accuracy captures classification correctness but fails to reflect model calibration error, latency under serving conditions, or gradient approximation error norms.

VI. CODE PROVENANCE AND REPRODUCIBILITY

VII. CODE PROVENANCE AND REPRODUCIBILITY

In compliance with academic integrity standards, all scripts, logs, and configuration matrices are maintained at

github.com/Scenic644/ML-Milestone_2. The components of our codebase are classified as follows:

- **Upstream Libraries (Reused as-is):** `galore-torch==1.0` for the core GaLore projector and gradient decomposition [1]; Hugging Face `peft` (v0.8.0) for standard LoRA adapter layers [2]; Hugging Face transformers for the roberta-base architecture.
- **Team-Authored Code:** `src/train.py` and execution orchestrators incorporating automated VRAM tracking via `torch.cuda.max_memory_allocated()`, GLUE dataset namespacing patches, and batch experimental logging to CSV. Minimal forward/backward verification routine in `src/sanity_check.py`.

VIII. CONCLUSION

This study successfully reproduces the empirical guarantees of Gradient Low-Rank Projection on GLUE fine-tuning. GaLore bridges the gap between memory efficiency and full-parameter learning, out-performing LoRA while approaching full-rank AdamW performance. The method proves to be an exceptionally viable training strategy for memory-constrained hardware, provided users account for the $\sim 60\%$ compute latency overhead induced by low-rank projections.

REFERENCES

- [1] J. Zhao, Z. Zhang, B. Chen, Z. Wang, A. Anandkumar, and Y. Tian, “GaLore: Memory-Efficient LLM Training by Gradient Low-Rank Projection,” in *Proceedings of the 41st International Conference on Machine Learning (ICML)*, Vienna, Austria, PMLR 235, 2024. Available: [arXiv:2403.03507](https://arxiv.org/abs/2403.03507)
- [2] E. J. Hu, Y. Shen, P. Wallis, Z. Allen-Zhu, Y. Li, S. Wang, L. Wang, and W. Chen, “LoRA: Low-Rank Adaptation of Large Language Models,” in *Proceedings of the 10th International Conference on Learning Representations (ICLR)*, 2022. Available: [arXiv:2106.09685](https://arxiv.org/abs/2106.09685)
- [3] I. Loshchilov and F. Hutter, “Decoupled Weight Decay Regularization,” in *Proceedings of the 7th International Conference on Learning Representations (ICLR)*, 2019. Available: [arXiv:1711.05101](https://arxiv.org/abs/1711.05101)
- [4] V. Lialin, N. Shivagunde, S. Muckatira, and A. Rumshisky, “ReLoRA: High-Rank Training Through Low-Rank Updates,” in *Proceedings of the 12th International Conference on Learning Representations (ICLR)*, 2024. Available: [arXiv:2307.05695](https://arxiv.org/abs/2307.05695)
- [5] Y. Liu, M. Ott, N. Goyal, J. Du, M. Joshi, D. Chen, O. Levy, M. Lewis, L. Zettlemoyer, and V. Stoyanov, “RoBERTa: A Robustly Optimized BERT Pretraining Approach,” *arXiv preprint arXiv:1907.11692*, 2019. Available: [arXiv:1907.11692](https://arxiv.org/abs/1907.11692)
- [6] A. Wang, A. Singh, J. Michael, F. Hill, O. Levy, and S. R. Bowman, “GLUE: A Multi-Task Benchmark and Analysis Platform for Natural Language Understanding,” in *Proceedings of the 7th International Conference on Learning Representations (ICLR)*, 2019. Available: [arXiv:1804.07461](https://arxiv.org/abs/1804.07461)